

Import the Libraries

In [1]:

```
import numpy as np
import pandas as pd
```

Dataset Loading

In [2]:

```
data=pd.read_csv(r"C:\Users\Acer\Desktop\Dataset\heart_disease dataset.csv")
```

In [3]:

```
data
```

Out[3]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	ta
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	
...	
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2	
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3	
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2	
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2	
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3	

1025 rows × 14 columns



In [4]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1025 non-null   int64  
 1   sex         1025 non-null   int64  
 2   cp          1025 non-null   int64  
 3   trestbps    1025 non-null   int64  
 4   chol        1025 non-null   int64  
 5   fbs         1025 non-null   int64  
 6   restecg     1025 non-null   int64  
 7   thalach     1025 non-null   int64  
 8   exang       1025 non-null   int64  
 9   oldpeak     1025 non-null   float64 
10   slope       1025 non-null   int64  
11   ca          1025 non-null   int64  
12   thal        1025 non-null   int64  
13   target      1025 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

In [5]:

```
data.isnull().sum()
```

Out[5]:

```
age         0
sex         0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
dtype: int64
```

In [6]:

```
data["target"].value_counts()
```

Out[6]:

```
1    526
0    499
Name: target, dtype: int64
```

Data Slicing

In [7]:

```
x=data.iloc[:, :-1].values  
y=data.iloc[:, -1].values
```

Using Imbalancing Technique to Balance data

In [8]:

```
from imblearn.over_sampling import SMOTE  
s=SMOTE()  
x_data,y_data=s.fit_resample(x,y)
```

In [9]:

```
from collections import Counter  
print("The Original data : ",Counter(y))  
print("The Aritificial data : ",Counter(y_data))
```

The Original data : Counter({1: 526, 0: 499})
The Aritificial data : Counter({0: 526, 1: 526})

Cross Validation Technique

In [10]:

```
from sklearn.model_selection import KFold  
kf=KFold(n_splits=6,random_state=5,shuffle=True)  
kf.get_n_splits(x_data)  
print(kf)
```

KFold(n_splits=6, random_state=5, shuffle=True)

In [11]:

```

for train_index, test_index in kf.split(x_data):
    print("TRAIN : ", train_index, "TEST : ", test_index)
    x_train, x_test = x_data[train_index], x_data[test_index]
    y_train, y_test = y_data[train_index], y_data[test_index]

    from sklearn.linear_model import LogisticRegression
    lr = LogisticRegression()
    from sklearn.model_selection import cross_val_score
    score = cross_val_score(lr, x_data, y_data, cv=kf)
    print(score)
    print(np.mean(score)*100)

```

```

TRAIN : [ 0  1  2  4  5  6  7  9 10 13 14 16 17
18
19 20 21 22 24 25 26 27 29 30 31 32 36 37
38 39 40 41 42 43 44 45 46 47 48 49 50 52
53 54 55 56 57 58 61 62 63 64 65 66 67 68
69 70 73 74 75 76 77 78 79 80 81 82 83 84
87 88 89 90 91 92 93 94 95 96 97 98 99 100
101 102 103 104 105 107 108 109 110 111 112 113 114 115
116 117 118 119 120 122 123 124 125 126 127 128 129 130
132 134 135 136 137 138 139 140 141 142 143 144 145 146
147 149 150 151 152 153 154 155 156 157 158 160 161 163
164 167 168 169 170 171 172 173 174 175 177 178 179 180
181 182 183 184 185 186 187 188 189 190 191 192 194 195
197 199 200 202 204 205 206 208 209 210 211 212 213 214
215 216 217 218 220 221 223 224 225 226 227 229 231 232
233 234 235 237 238 239 240 241 242 243 245 246 247 248
249 251 252 253 255 256 257 259 260 261 262 263 264 265
266 268 269 270 271 272 273 274 275 276 277 278 280 281
282 284 285 287 288 289 290 291 292 294 295 296 297 298
299 301 302 303 304 305 306 307 308 310 311 312 314 315

```

Predicting the Test set results

In [12]:

```
from sklearn.model_selection import cross_val_predict
y_pred=cross_val_predict(lr,x_test,y_test,cv=kf)
print(y_pred)
```

C:\Users\Acer\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
C:\Users\Acer\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:
444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
C:\Users\Acer\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:
444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(

[1 0 1 0 0 0 1 0 0 1 1 1 1 0 0 1 1 1 1 1 1 0 1 1 1 1 0 1 0 1 0 1 0 0 0 1 1 1
 1 1 0 0 1 1 0 1 1 1 1 1 1 1 1 0 0 1 1 0 1 0 0 1 1 0 1 1 1 0 0 0 0 1 1 0 0
 0 1 0 1 1 1 0 1 1 0 1 1 1 0 0 1 1 1 1 0 0 1 0 1 0 0 0 1 1 0 1 1 1 0 1 1 0
 0 0 1 0 1 0 0 0 1 0 1 1 1 1 1 1 1 1 1 0 1 0 1 1 1 1 0 0 1 0 1 0 1 0 1 1 1
 1 0 1 0 1 1 1 0 1 1 1 1 1 0 0 0 0 0 1 1 0 0 0 0 1 1 0]
```

C:\Users\Acer\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\Acer\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:
444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
C:\Users\Acer\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:
444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Confusion Matrix , Accuracy score and Classification Report

In [13]:

```

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
ac = accuracy_score(y_pred, y_test) * 100
cm = confusion_matrix(y_pred, y_test)
cr = classification_report(y_pred, y_test)
print("Accuracy score:", ac)
print("*****")

print("Confusion matrix:")
print(cm)
print("*****")

print("classification report:")
print(cr)

```

Accuracy score: 79.42857142857143

Confusion matrix:

[[57 14]

[22 82]]

classification report:

	precision	recall	f1-score	support
0	0.72	0.80	0.76	71
1	0.85	0.79	0.82	104
accuracy			0.79	175
macro avg	0.79	0.80	0.79	175
weighted avg	0.80	0.79	0.80	175

In []: